K. S. Sushanth

+91 63791 19460 | k.s.sushanth06@gmail.com | linkedin.com/in/sushanthks/ | github.com/SushanthKS06

Education

Bachelor of Technology in Information Technology — Sep 2023 – May 2027 (Expected)
Sri Krishna College Of Engineering And Technology, Coimbatore, India — *CGPA: 8.11/10*

Technical Skills

Programming: Python, SQL, Bash, Java, Scala, Golang
Machine Learning: Regression, Classification, Clustering, Feature Engineering, Model Evaluation, NLP, Deep Learning
GenAI: Large Language Models (LLMs), Prompt Engineering, RAG, Embeddings
Computer Vision: OpenCV, YOLOv8
Frameworks & Tools: Scikit-learn, PyTorch, Hugging Face, LangChain, FastAPI, Streamlit
Cloud & Infra: AWS (SageMaker/Lambda), GCP, Docker, Kubernetes, Git, Github, Linux
Database: MySQL, PostgreSQL, MongoDB, Redis, Cassandra, ChromaDB, Pinecone, FAISS
Data & Analytics: Pandas, NumPy, EDA, SQL Optimization, PySpark, Databricks, Hadoop

Experience

AI Engineer Intern — Aug 2025 – Nov 2025
Infosys Springboard – Applied AI Projects — *Remote*

- Built ML and GenAI systems for applied use cases where downstream users depended on outputs to proceed, and incomplete requirements were clarified only after failures surfaced
- Prevented wasted experimentation by establishing classical ML baselines upfront to decide whether higher-complexity models were justified, reducing iteration time by **30%**
- Developed GenAI pipelines for document summarization and information extraction; early hallucinated outputs caused manual rework and blocked adoption, leading to schema-constrained generation that reduced review effort by **60%**
- Detected silent data quality issues only after model outputs degraded across repeated runs; introduced validation and monitoring checks to prevent regressions from reaching users
- Worked directly with technical reviewers and non-technical users to define accuracy, latency, and reliability thresholds that determined whether outputs were usable or rejected

Data Analyst Intern — Jun 2025 – July 2025
Elevate Labs — *Remote*

- Analyzed **100k+** records used by internal teams to make weekly prioritization and reporting decisions
- Identified SQL queries that failed under higher data volumes and blocked report refreshes; refactored multi-join queries to improve execution time by **40%**
- Automated recurring analytics workflows using Python and SQL, freeing **10+ hours/week** and allowing teams to act on updated metrics earlier
- Converted vague stakeholder questions into concrete KPIs after earlier reports failed to drive decisions

Projects

Auto-Vision: Vehicle Damage Assessment System | *Python, YOLOv8, OpenCV, FastAPI, Docker*

- Built a vehicle inspection pipeline to detect and classify scratches and dents from automobile images, designed around marketplace inspection constraints
- Initial models failed under low-light conditions common in real uploads, causing unreliable assessments; iterated on data augmentation and model variants to achieve **92%** detection accuracy
- Designed a Random Forest–based pricing adjustment model to estimate damage severity, ensuring rare defect classes did not disproportionately skew pricing
- Optimized inference latency to **sub-200ms** so inspections could proceed in real time without slowing evaluation workflows

Smart-Doc: GenAI Document Intelligence System | *Python, Llama-3, LangChain, FAISS*

- Built a retrieval-augmented GenAI system to extract structured fields from legal and financial documents used in downstream automation
- Early outputs contained fabricated fields that blocked automated processing; reworked prompts and validation logic to produce reliable JSON suitable for SQL ingestion
- Selected FAISS over managed vector stores to reduce latency and avoid recurring infrastructure costs under large-context workloads

Market Pulse: Predictive Sales Analytics Engine | *Python, SQL, Pandas, XGBoost, Streamlit*

- Processed and cleaned **500,000+** transactional records after identifying missing values and outliers that distorted early forecasts
- Trained and tuned regression models achieving **MAE ¡ 5%**, replacing weaker baseline forecasts previously used for planning
- Delivered dashboards that directly changed inventory prioritization and replenishment decisions

Certifications & Achievements

Certifications: AWS Certified Developer – Associate; DeepLearning.AI (Generative AI with LLMs, Multi-Agent Systems); Hugging Face (Deep Reinforcement Learning); Google Cloud Fundamentals
Problem Solving: Solved **250+** algorithmic problems on LeetCode and Codeforces